

COMP6237 Data Mining

Lecture 10: Semantic Spaces (Finding Features I)

Zhiwu Huang

Zhiwu.Huang@soton.ac.uk

Lecturer (Assistant Professor) @ VLC of ECS
University of Southampton

Lecture slides available here:

<http://comp6237.ecs.soton.ac.uk/zh.html>

(Thanks to Prof. Jonathon Hare and Dr. Jo Grundy for providing the lecture materials used to develop the slides.)

Semantic Spaces – Problem

	doc	title
0	m1	Human machine interface for ABC computer applications
1	m2	A survey of user opinion of computer system response time
2	m3	The EPS user interface management system
3	m4	System and human system engineering testing of EPS
4	m5	Relation of user perceived response time to error measurement
5	g1	The generation of random, binary, ordered trees
6	g2	The intersection graph of paths in trees
7	g3	Graph minors IV: Widths of trees and well-quasi-ordering
8	g4	Graph minors: A survey

Are these documents about the same topic?

Semantic Spaces – Problem

Semantic Spaces:

- ▶ represent word meanings as vectors that keep track of the words distributional history
- ▶ focus on semantic similarity
- ▶ similarity measured using geometrical methods

e.g. Cosine similarity between **PC** and **Windows** = 0.77

Cosine similarity between **PC** and **window** = 0.13

In Japanese, A. Utsumi, IEEE SMC 2010

Semantic Spaces – Latent Semantic Analysis


Consider a term-document matrix which describes occurrences of terms in documents [Demo: Semantic_Spaces_Visual](#)

	m1	m2	m3	m4	m5	g1	g2	g3	g4
computer	0.58	0.46	0.00	0.00	0.00	0.0	0.00	0.00	0.00
eps	0.00	0.00	0.60	0.71	0.00	0.0	0.00	0.00	0.00
graph	0.00	0.00	0.00	0.00	0.00	0.0	0.71	0.55	0.52
human	0.58	0.00	0.00	0.71	0.00	0.0	0.00	0.00	0.00
interface	0.58	0.00	0.60	0.00	0.00	0.0	0.00	0.00	0.00
minors	0.00	0.00	0.00	0.00	0.00	0.0	0.00	0.63	0.60
perceived	0.00	0.00	0.00	0.00	0.58	0.0	0.00	0.00	0.00
response	0.00	0.46	0.00	0.00	0.49	0.0	0.00	0.00	0.00
survey	0.00	0.46	0.00	0.00	0.00	0.0	0.00	0.00	0.60
time	0.00	0.46	0.00	0.00	0.49	0.0	0.00	0.00	0.00
trees	0.00	0.00	0.00	0.00	0.00	1.0	0.71	0.55	0.00
user	0.00	0.40	0.52	0.00	0.43	0.0	0.00	0.00	0.00


Semantic Spaces – Latent Semantic Analysis

Latent Semantic Analysis (LSA) makes a low-rank approximation
It assumes the term-document matrix:

- ▶ is noisy, and should be de-noised
- ▶ is more sparse than it should be



Semantic Spaces – Recap SVD

$$A = U \Sigma V^T$$



The diagram shows the SVD decomposition $A = U \Sigma V^T$. Matrix A is $m \times n$. Matrix U is $m \times p$. Matrix Σ is $p \times p$. Matrix V^T is $p \times n$. The dimensions are indicated in blue text below each matrix name.

Where p is rank of matrix A

U called *left singular vectors*, contains the eigenvectors of AA^T ,
 V called *right singular vectors*, contains the eigenvectors of $A^T A$
 Σ contains square roots of eigenvalues of AA^T and $A^T A$

If A is matrix of mean centred feature vectors, V contains principal components of the covariance matrix

Semantic Spaces – Recap Truncated SVD




Uses only the largest r singular values (and corresponding left and right vectors)

This can give a *low rank approximation* of A , $\tilde{A} = U_r \Sigma_r V_r$

This has the effect of minimising the Frobenius norm of the difference between A and $\tilde{A}$

Semantic Spaces – Recap Truncated SVD




Each row of V_r corresponds to an eigenvector of $A^T A$

- ▶ This means it is proportional to the covariance or correlation between the documents
- ▶ These are the *concepts*

Each row of U_r describes a term as a vector of weights with respect to r *concepts*

Each column of V_r describes a document as a vector of weights with respect to r *concepts*


Semantic Spaces – LSA



Term concepts and document concepts have the same dimensionality, but represent different spaces.

Semantic Spaces – LSA


$$\text{SVD: } A = U\Sigma V^T$$



We then reduce the dimensionality by choosing only the first few eigenvalues (in Σ) and the corresponding columns in U and V .

Semantic Spaces – LSA


$$\text{SVD: } A \approx U_r \Sigma_r V_r^T \quad r = 4$$



Each row of U_r describes a word as a vector of weights with respect to r concepts

Each column of V_r describes the title as a vector of weights with respect to r concepts

Semantic Spaces – LSA




Semantic Spaces – LSA

What do these 'concepts' mean?

U_r gives us the weighting of the words for each concept

We can show that weighting for each concept:




This shows us perhaps that the concepts are not always particularly meaningful.

[Demo: Semantic_Spaces_Visual](#)

Semantic Spaces – LSA

Cosine similarity of document vectors can be compared ($r = 2$)

- ▶ Vectors for “m1” and “m2” give cosine similarity = 0.93
- ▶ Vectors for “g1” and “g2” give cosine similarity = 0.83
- ▶ Vectors for “g1” and “m1” give cosine similarity = 0.18



```
r = 2
Ur = U[:, :r]
Sigmar = np.diag(s[:r])
VTr = VT[:, r, :]


doc_coords = (Sigmar @ VTr).T
word_coords = Ur @ Sigmar
```

Clustering algorithms can be used on the vectors

Semantic Spaces – LSA

Cosine similarity of word vectors can be compared ($r = 2$)

- ▶ Vectors for “human” and “interface” give cos similarity = 0.95
- ▶ Vectors for “human” and “user” give cos similarity = 0.11
- ▶ Vectors for “graph” and “minor” give cos similarity = 0.90



```
r = 2
Ur = U[:, :r]
Sigmar = np.diag(s[:r])
VTr = VT[:r, :]

doc_coords = (Sigmar @ VTr).T
word_coords = Ur @ Sigmar
```

Clustering algorithms can be used on the vectors

Semantic Spaces – LSI


Latent Semantic Indexing (LSI) LSA can be used for document retrieval

- ▶ Given Query: view as query vector $\mathbf{q}$
- ▶ Project $\mathbf{q}$ in to document space
- ▶ Compare with document vectors, find closest

Results work mathematically

However, results may not be easy to interpret in terms of natural language.

Semantic Spaces – LSI



Semantic Spaces – LSA

Problems?

- ▶ Polysemious words - with multiple meanings - aren't captured
 - ▶ The vector representation averages all meanings of the word
 - ▶ e.g. 'fit' is an adjective and a verb
- ▶ Word order is ignored (use n-grams?) *An n-gram is a sequence of n words*
- ▶ LSA assumes words and documents form a joint Gaussian distribution, however a Poisson distribution is observed

[Demo: Semantic_Spaces_Visual](#)


Semantic Spaces – LSA

So far: *bag of words* (BOW) from natural language.

However the maths should work for compositions of occurrences in any unit.

For example, in image search we might want to search for other images with circles.

The image could be encoded with the number of different shapes it has.



BoW: ○ △ □ × *

 2 0 3 0 0

Each image becomes a bag of visual words

img_A	2	3	0	0
img_B	0	0	5	0
img_C	1	2	0	0
	circle	square	triangle	star

Semantic Spaces – LSA


Need to make a large multidimensional space in which images, keywords and visual terms can be placed

In training:

- ▶ Learn how images and keywords are related
- ▶ Place images and keywords close together in the space

Unannotated images can be placed in the space based on the visual terms they contain

- ▶ Images can be placed based on their visual terms in the space
- ▶ They should lie near the keywords that describe them




[Demo: Semantic_Spaces_Visual](#)

Semantic Spaces – LSA

This lower dimensional space can be used to:

- ▶ Find Images using similar words
- ▶ Find images with similar images
- ▶ Return possible key words for an image
- ▶ Find relationships between words, and between words and visual terms



Semantic Spaces – Summary


Top query retrieval results for: 'user response time'

	doc	title	cosine_with_query
4	m5	Relation of user perceived response time to error measurement	0.999965
0	m1	Human machine interface for ABC computer applications	0.999252
2	m3	The EPS user interface management system	0.999186
3	m4	System and human system engineering testing of EPS	0.998190
1	m2	A survey of user opinion of computer system response time	0.994227

Takeaway:

- TF-IDF gives a sparse term-document matrix
- SVD / truncated SVD creates a lower-dimensional latent space
- Cosine similarity in that space captures semantic closeness
- The same idea extends beyond text to visual words and multimodal data

Semantic Spaces – Roadmap



Finding Features I – Textbook

CHAPTER 10

Finding Independent Features

Most of the chapters so far have focused primarily on *supervised* classifiers, except Chapter 3, which was about *unsupervised* techniques called *clustering*. This chapter will look at ways to extract the important underlying features from sets of data that are not labeled with specific outcomes. Like clustering, these methods do not seek to make predictions as much as they try to characterize the data and tell you interesting things about it.

- ▶ Programming Collective Intelligence: Building Smart Web 2.0 Applications *T. Segaran*.

Semantic Spaces – Learning Outcomes

- **LO1:** Demonstrate an understanding of techniques for finding independent semantic features, such as: (exam)
 - ❖ Comprehending the core concepts of Latent Semantic Analysis (LSA) and apply LSA on a dataset
 - ❖ Discussing the advantages and disadvantages of algorithm
- **LO2:** Implement the learned algorithms for independent semantic feature learning (coursework)

Assessment hints: Multi-choice Questions (single answer: concepts, calculation etc)

- *Textbook Exercises: textbooks (Programming + Mining)*
- *Other Exercises: <https://www-users.cse.umn.edu/~kumar001/dmbook/sol.pdf>*
- *ChatGPT or other AI-based techs*