

# COMP6237 – Linear Regression

Shoaib Ehsan

[s.ehsan@soton.ac.uk](mailto:s.ehsan@soton.ac.uk)

Lecture slides available here:

<http://comp6237.ecs.soton.ac.uk/>

(Credit goes to Jon Hare who developed a large part of the module)

# General Plan

---

- At the start of the course:
  - An introduction to regression techniques
- These lectures are based on the stats package R
  - This is free software, you can download it from <https://www.r-project.org/>
  - If you are not familiar with R, follow a tutorial to get some idea:  
<https://www.southampton.ac.uk/~mb1a10/Rtutorial/R.html>
- At the end of the course:
  - An introduction to Information Theory
  - Mining Data Streams

# COMP6237: Linear Regression

---

## .Outline:

- Brief revision of some basic stats
- Variables and prediction
- The method of least squares (LS)
- Practical implementation in R
- Linear regression in higher dimensions
- Maximum likelihood estimation (MLE)
- LS and MLE
- Weighted LS, Heteroskedasticity, and local linear regression
- Problem sheet for linear regression and MLE:  
– <https://www.southampton.ac.uk/~mb1a10/stats/exercise1.pdf>

# Reminder of Some Basic Stats (1)

---

- .Suppose we have a set of  $N$  observations/measurements
- .Can analyse them via histograms/pdfs  $\{X_1, X_2, \dots, X_N\}$
- .How to classify distributions?
  - Means (Median, mode, etc.) – measures of central tendency for a distribution

$$E[X] = 1/N \sum_{i=1}^N X_i$$

- Variances/standard deviation – measures of dispersion

$$V[X] = 1/N \sum_{i=1}^N (X_i - E[X])^2 = E[X^2] - E[X]^2$$

- Could use the Central Limit Theorem to argue about standard errors, confidence intervals, etc.

# Reminder of Some Basic Stats (2)

---

• We are often interested in relationships between pairs of observations

$$\{(X_1, Y_1), (X_2, Y_2), \dots, (X_N, Y_N)\}$$

• Can classify relationships in various ways

- Covariance (measure of the joint variability of two random variables)

$$\text{Cov}[X, Y] = 1/N \sum_{i=1}^N (X_i - E[X]) (Y_i - E[Y])$$

- Correlation coefficients, e.g. Pearson

$$r = \frac{\text{Cov}[X, Y]}{\sqrt{V[X]} \sqrt{V[Y]}}$$

- $R^2$  ... proportion of variance of X explained by knowledge of Y

# What about Prediction?

---

- What about if we want to predict the value of one variable based on the knowledge about another variable?
- This is what regression analysis is all about

# Interlude: Types of Variables (1)

---

.Three types of variables:

- **Continuous**: real numbered values (e.g., time, mass)
- **Ordinal**: a numerical variable where a small number of possibilities are ranked (e.g., school grades)
  - . Boundaries between first two types can be blurred (e.g. most “continuous” variables are really ordinal because they can only be measured up to some accuracy)
- **Categorical**: describes membership of a group (but cannot be ranked). e.g., country of birth, gender, etc.
  - . Some are binary and others are not

# Interlude: Types of Variables (2)

---

.Whatever form a variable has, they can play different roles when we build statistical models

## **.Dependent or outcome variables**

- For our analysis this variable is the focus. It is assumed to be predictable from some other variables.

## **.Independent or predictor variables**

- Are assumed to have inherent variation (“they just are”). We will use them to explain the variance in the dependent variables.

# Example: Demographic Factors

---

.Let's assume we want to predict driving test outcomes from demographic data.

.For this particular analysis:

- Independent variables: demographic factors like gender, age, weight, height, etc.
- Dependent variables: driving test outcomes like outcome of the practical/theoretical parts


.Role of variables depends on analysis/model we have in mind.

- We could equally well be interested in the inverse, i.e., predicting demographic factors based on knowledge of driving test outcomes.

# Variables in Regression

---

- Simplest case involves one dependent (Y) and one independent (X) variable
- Both variables are continuous (or at least ordinal)
- We are trying to predict the variation in Y based on the variation in X




Given a plot like this:  
What are your options?

Easiest way:  
Assume linear relationship  
between X and Y

Try to find line of “best guess”

# Drawing a Line Through Some Points


---



- Look for line of “best fit” through cloud of X,Y points
- $Y=mX+b$ , Need to find two parameters  $m$  and  $b$
- How to be systematic about it?

# Method of Least Squares

- For any given line  $mX+b$  we can measure the difference between our actual  $Y$  values and those predicted
- Squared differences** are a reasonable measure of the goodness of fit; regression analysis tries to minimize this difference



# LS as an Optimisation Problem

---

.One might think that minimizing the squared differences is a difficult combinatorial optimisation problem ... it is not:

$$E(m, b) = \sum_{i=1}^N (mX_i + b - Y_i)^2 = \min!$$

$$\frac{\partial}{\partial b} E(m, b) = 2 \sum_{i=1}^N (mX_i + b - Y_i) = 0$$

$$b = E[Y] - mE[X]$$

.Substituting back:

$$E(m, b) = \sum_{i=1}^N (m(X_i - E[X]) - (Y_i - E[Y]))^2$$

$$\frac{\partial}{\partial m} E(m, b) = 2 \sum_{i=1}^N (m(X_i - E[X]) - (Y_i - E[Y])) (X_i - E[X]) = 0$$

$$m = \frac{\sum_{i=1}^N (Y_i - E[Y]) (X_i - E[X])}{\sum_{i=1}^N (X_i - E[X])^2} = \text{Cov}[X, Y] / V[X]$$

# Regression in Machine Learning


---

- Slightly different point of view is that we might consider the pairs  $(X_i, Y_i)$  as a training set
- Want to learn a mapping  $y=f(x)$  from the training set
- Approach often:
  - “somehow” parametrise  $f(x)$  (e.g.,  $f(x)=mx+b$  here) and try to learn best parameters  $m$  and  $b$
  - This is often done via minimising some error function  $E$  which sums up squared residual errors over training set (i.e., this is the same as above)
  - e.g., in NN  $f(x)$  is a nonlinear function

# How Do I Run Linear Regression in R?

---

- Let's start with a fictional data set
  - Say a 100 men have been measured for their height and basketball ability
  - We want to predict basketball skill from height



Mean height is 169cm, stddev 10.5cm  
Mean ability is 66 (arbitrary scale), stdev 6

Correlation  $r=0.52$  (i.e. 28% of the variation of basketball ability are explained by height)

# R Commands ...

---

- Read data into R with the usual command `read.table ()`; the variables of interest are Height and BasketballAbility
- Build a regression model with
- `regmodel=lm (BasketballAbility ~ Height)`
- (lm stands for “linear model”)
- `summary (regmodel)` gives us most of the information we need ...

# The Full Output

Residuals:

| Min      | 1Q      | Median  | 3Q     | Max     |
|----------|---------|---------|--------|---------|
| -11.0733 | -3.4851 | -0.5733 | 3.4969 | 12.9267 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t ) |     |
|-------------|----------|------------|---------|----------|-----|
| (Intercept) | 15.67832 | 8.22351    | 1.907   | 0.0595   | .   |
| Height      | 0.29602  | 0.04855    | 6.097   | 2.15e-08 | *** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.07 on 98 degrees of freedom

Multiple R-squared: 0.275, Adjusted R-squared: 0.2676

F-statistic: 37.17 on 1 and 98 DF, p-value: 2.146e-08

# The Full Output

Residuals:

| Min      | 1Q      | Median  | 3Q     | Max     |
|----------|---------|---------|--------|---------|
| -11.0733 | -3.4851 | -0.5733 | 3.4969 | 12.9267 |

Characterizes the distribution of “residuals” (differences between predicted output and actual output)

Should roughly be normally distributed with a mean of zero

# The Full Output

This is the important test here, i.e. the hypothesis of a relationship (slope  $\neq 0$ ) is confirmed.

Residuals:

| Min      | 1Q      | Median  | 3Q     | Max     |
|----------|---------|---------|--------|---------|
| -11.0733 | -3.4851 | -0.5733 | 3.4969 | 12.9267 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t ) |     |
|-------------|----------|------------|---------|----------|-----|
| (Intercept) | 15.67832 | 8.22351    | 1.907   | 0.0595   | .   |
| Height      | 0.29602  | 0.04855    | 6.097   | 2.15e-08 | *** |

The coefficients (“m” and “b”).

Meaning ... a man of zero height would have basketball ability 15.7,  
Every additional cm of height adds 0.3 to basketball ability.

Also get standard errors, t-tests for hypothesis that values are 0

# Generalisations

- Multi-dimensional input?

- “fit” a hyperplane

$$\mathbf{x} \in \mathbb{R}^D, y \in \mathbb{R}$$

- A general linear regression function then is

$$f(\vec{x}) = \sum_{j=1}^D w_j x_j + b = \mathbf{w}\mathbf{x} + b$$

- For simpler notation we could say

$$\tilde{\mathbf{w}} = [w_1, w_2, \dots, w_D, b]^T, \tilde{\mathbf{x}} = [x_1, x_2, \dots, x_D, 1]^T$$

- Hence:

$$E = \sum_{i=1}^N (y_i - \tilde{\mathbf{w}}\tilde{\mathbf{x}}_i)^2$$




Figure 3.1: Linear least squares fitting with  $X \in \mathbb{R}^2$ . We seek the linear function of  $X$  that minimizes the sum of squared residuals from  $Y$ .

# Generalisations (2)

---

- For a more compact notation say

$$\tilde{X} = \begin{bmatrix} \tilde{\mathbf{x}}_1^T & . \\ \tilde{\mathbf{x}}_2^T & . \\ . & . \\ . & . \\ \tilde{\mathbf{x}}_N^T & . \end{bmatrix} \quad \mathbf{y} = [y_1, y_2, \dots, y_N]^T$$

- Then:  $E = \sum_{i=1}^N (y_i - \tilde{\mathbf{w}} \tilde{\mathbf{x}}_i)^2 = \|\mathbf{y} - \tilde{X} \tilde{\mathbf{w}}\|^2 = (\mathbf{y} - \tilde{X} \tilde{\mathbf{w}})^T (\mathbf{y} - \tilde{X} \tilde{\mathbf{w}})$
- Can find  $\mathbf{w}$  through  $dE/d\mathbf{w}_i = 0 \dots$

$$\tilde{\mathbf{w}} = (\tilde{X}^T \tilde{X})^{-1} \tilde{X}^T \mathbf{y} = X^+ \mathbf{y}$$



pseudoinverse of  $X$  also denoted by  $X^+$

# Generalisations (3)

---

- Some practical considerations

- The data matrix  $X$  can have large dimensions and we might need to calculate an inverse ... this can be quite time consuming

- Two ways to go about it:

- Often via **QR factorization**, factorizing  $\tilde{X} = QR$
- with  $Q$  a diagonal matrix and  $R$  an upper triangular matrix (can be done via applying Gram-Schmidt procedure to column vectors) → for exact solutions
- Alternatively: can use **stochastic gradient descent** to numerically minimize residuals

# Generalisations (4)

---

- D dimensional input, K dimensional output; could say:

$$\mathbf{y} = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}} \quad \tilde{\mathbf{W}} = [\tilde{\mathbf{w}}_1, \tilde{\mathbf{w}}_2, \dots, \tilde{\mathbf{w}}_K] = \begin{bmatrix} \mathbf{w}_1 & \mathbf{w}_2 & \dots & \mathbf{w}_K \\ b_1 & b_2 & \dots & b_K \end{bmatrix}$$

This is useful since:  $y_j = \tilde{\mathbf{w}}_j^T \tilde{\mathbf{x}}$

- An error function can be constructed by summing over all pairs and output dimensions

$$E = \sum_{i=1}^N \sum_{j=1}^K (y_{i,j} - \tilde{\mathbf{w}}_j^T \tilde{\mathbf{x}}_i)^2$$

- Introduce:  $\mathbf{y}'_j = [y_{1,j}, y_{2,j}, \dots, y_{N,j}]^T$   $\tilde{\mathbf{X}}^T = [\tilde{\mathbf{x}}_1, \tilde{\mathbf{x}}_2, \dots, \tilde{\mathbf{x}}_N]$

$$\rightarrow E = \sum_{j=1}^K \|\mathbf{y}'_j - \tilde{\mathbf{X}} \tilde{\mathbf{w}}_j\|^2$$

- i.e. we have K distinct estimation problems with solutions  $\tilde{\mathbf{w}}_j = \tilde{\mathbf{X}}^+ \mathbf{y}'_j$

# Some Comments

---

.One may have wondered if different results would have been obtained when making a different choice about the difference function between predictions and data points?

- i.e., what if one wouldn't use the sum of the squares ( $\sim L_2$  norm) but some other measure?  
→ other results would have been obtained!

.Least squares is a very common strategy in the sciences

.Another prominent approach is maximum likelihood estimates

# Maximum Likelihood Estimation

---

- Suppose we have a set of  $n$  data points  $X_1, \dots, X_n$  which are from some pdf  $f(x;p)$  which has some “hidden” parameters  $p$ . We want to “guess” these parameters.
- How to go about it? Construct a likelihood function:

$$L(X_1, X_2, \dots, X_n; p) = \prod_{i=1}^n f(X_i; p)$$

Likelihood of obtaining the data by sampling from  $f$  given the parameter  $p$

- Maximising  $L$  will give us a value of  $p$  corresponding to the most likely explanation of the data

# Example (1)

---

• Suppose our observations have been generated from a normal distribution with unknown mean and variance, i.e.,

$$X \sim N(\mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(\frac{-(x - \mu)^2}{2\sigma^2}\right)$$

• Then the likelihood function is:

$$L(X_1, \dots, X_n; \mu, \sigma^2) = \sigma^{-n} (2\pi)^{-n/2} \exp\left(\frac{-1}{2\sigma^2} ((X_1 - \mu)^2 + (X_2 - \mu)^2 + \dots + (X_n - \mu)^2)\right)$$

• L is max if  $\sum_{i=1}^N (X_i - \mu)^2 = \min!$       which is if  $\mu = 1/N \sum_{i=1}^N X_i$

• Hence the expectation is a **maximum likelihood estimator** for this example.

# Example (2)

---

.Suppose we have  $n$  observations  $X_1, \dots, X_n$  that have been sampled from the uniform distribution over  $[0, N]$  and  $N$  is unknown.

# Example (2)

---

.Suppose we have  $n$  observations  $X_1, \dots, X_n$  that have been sampled from the uniform distribution over  $[0, N]$  and  $N$  is unknown.

.Construct likelihood function

$$L(X_1, \dots, X_n; N) = \begin{cases} 0 & \text{any } X_i \text{ outside } [0, N] \\ (1/N)^n & \text{otherwise} \end{cases}$$

.Maximum likelihood estimator?

## Example (2)

---

• Suppose we have  $n$  observations  $X_1, \dots, X_n$  that have been sampled from the uniform distribution over  $[0, N]$  and  $N$  is unknown.


• Construct likelihood function

$$L(X_1, \dots, X_n; N) = \begin{cases} 0 & \text{any } X_i \text{ outside } [0, N] \\ (1/N)^n & \text{otherwise} \end{cases}$$

• Maximum likelihood estimator is  $N = \max(X_1, X_2, \dots, X_n)$

# Implied Model of Data Generation in Linear Regression

---



# MLE and LS

---

• Suppose we know a priori that  $X$  and  $Y$  have a linear relationship except for some noise, i.e.,

$$Y_i = mX_i + b + \epsilon_i, \epsilon_i \sim N(0, \sigma^2)$$

• where the epsilons are independent. We aim to find  $m$  and  $b$  via MLE:

$$L(X_1, \dots, X_n; m, b) = (2\pi)^{-n/2} \sigma^{-n} \exp\left(\frac{-1}{2\sigma^2} \sum_i (Y_i - mX_i - b)^2\right)$$

• To maximise  $L$  we need to minimise the sum of squares in the exponent, i.e.,

- Least squares is equivalent to an MLE estimate for  $m$  and  $b$  if  $X$  and  $Y$  are **linearly** related with **Gaussian noise**. Sum of squares has a privileged position ...

# Weighted Least Squares

---

• Instead of minimising

$$\sum_{i=1}^N (mX_i + b - Y_i)^2 = \min!$$

• one might want to minimise:

$$\sum_{i=1}^N w_i (mX_i + b - Y_i)^2 = \min!$$

• Why?

- To focus accuracy – might be more interested in certain X-regions, or errors in some regions might be more costly than in others
- There is a number of other optimisation problems transformed/approximated by WLS, e.g., generalised linear models where the response is some nonlinear function of a linear predictor (e.g., logistic regression, see later)

# Homo-/Heteroskedasticity


---

## .Discounting imprecision


- Ordinary least squares assumes  $y=mx+b+e$  with  $e$  iid Gaussian white noise. This implies that  $e$  has constant variance (**homoskedasticity**).
- Often this is not the case → **heteroskedastic data**
- Can then set  $w_i=1/s_i^2$  so we get the heteroskedastic MLE
- (in other words: does not make much sense to concentrate on noisy parts of the data, want to use parts with little noise for our estimates)

# Homo-/Heteroskedasticity – Examples

---





Homoscedasticity



Heteroscedasticity

# An Example in R

.Suppose we have a linear relationship  $y=3-2x$  between  $X$  and  $Y$ , and add Gaussian white noise with  $s(x)=1+0.5x^2$




```
x=rnorm(100,0,3)
Y=3-2*x+rnorm(100, 0, sqrt(1+0.5*x^2))
plot(x,y)
abline(a=3,b=-2,col="grey")
fit.ols= lm(y~x)
abline(fit.ols$coefficients,lty=2)
```

plot(x,fit.ols\$residuals)

As expected, variance is not constant!  
Fit misses the real relationship by some margin.

# Weighted Linear Regression



unweighted regression

weighted regression performs better

real relationship

```
fit.wls = lm (y~x, weights=1/(1+0.5*x*x))  
abline(fit.wls$coefficients, lty=3)
```


# How to know the proper weights?

---

- Somehow know it from measurement device (e.g., know precision of the device for various ranges)
- e.g., in polls or surveys variance of the proportions we find should be inversely related to sample size, hence can make weights proportional to sample size.
- Try to estimate it from the data, e.g.
  - Estimate  $y(x)$
  - Construct log squared residuals  $z_i = \log((y_i - r(x_i))^2)$
  - Estimate mean of the  $z$ 's  $\rightarrow q(x)$
  - Use  $s_x^2 = \exp q(x)$


# What if the data are not linear?

---



# Local linear regression

---



# Local Linear Regression

---

- Linear regression could be justified by looking at a general regression function  $r(x)$  and expanding into a Taylor series

$$r(x) = r(x_0) + (x - x_0)r' + 1/2(x - x_0)^2 r''(x_0) + \dots (*)$$

- Then as long as we are close to  $x_0$   $r'$  is the regression coefficient, but what if the relationship is not linear?
- Naive approach: use some window (say of size  $h$ ) around data points and set

$$w_i = \begin{cases} 1 & \text{if } |x_i - x_0| < h \\ 0 & \text{otherwise} \end{cases}$$

... and then use weighted least squares

# Local Linear Regression (2)

---

.Often one wants weights to change a bit more smoothly than that.

.Kernel regression:

– Cut off Taylor expansion (\*) after constant term and solve

$$\min_b \sum_{i=1}^N w_i(x) (y_i - b)^2$$

– Which is solved by


$$b = \frac{\sum_{i=1}^N w_i(x) y_i}{\sum_{i=1}^N w_i(x)}$$

$$w_i(x) \propto K(x_i, x)$$

# Locally Linear Regression (3)

---

- Take 0<sup>st</sup> and 1<sup>st</sup> order terms from (\*)
- Often use a tri-cubic kernel




$$K(x_i, x) = \left(1 - \left(\frac{|x - x_0|}{h}\right)^3\right)^3$$

(h=1)

- R functions: lowess (specify fraction f of data points included), loess

# Example



```
x.sin=runif(100,0,1)
y.sin=sin(10*x.sin)+rnorm(100,0,1)
plot(x.sin,y.sin, xlab="x", ylab="y")
curve(sin(10*x),col="grey",add="TRUE")
lines(lowess(x.sin, y.sin,f=1/3),lty=1)
lines(lowess(x.sin, y.sin,f=2/3),lty=2)
lines(lowess(x.sin, y.sin,f=1/6),lty=3)
```

Some art in choosing  $f$  appropriately; but can do quite a good job.

# Local Linear Regression (4)

---

.When would you actually use this?

- Number of predictors is small
- You don't want to think too hard about what features to use
- (we will talk about other ways to deal with non-linear data later)

. Cons:

- Linear regression is a parametric technique: estimate weights from data and can then use to predict
- This is a non-parametric technique (the “data provide the function” – basically need to keep data points in memory to make predictions)
- Need more data ...

# Summary

---

## •Linear regression

- What it is, how it works
- How to judge significance
- Generalisations to higher d's

## •MLE

- Formalisation
- Relationship to LR

## •Extensions of linear regression

- Homo vs heteroskedastic data
- Weighted linear regression

# Problem (1)

---

A data set is constructed by taking 100 samples from a normal distribution with mean 5 and standard deviation 2 to construct a variable  $X_i$ ,  $i=1, \dots, 100$ . Then, a variable  $Y_i$ ,  $i=1, \dots, 100$  is constructed by taking the values of the corresponding  $X_i$  and adding one half of a random variate drawn from a normal distribution with mean 5 and standard deviation 2 and thus a set of 100 pairs  $(X_i, Y_i)$  is obtained.

Find the parameters of a linear regression of  $Y$  on  $X$ .

# Problem (2)

---

Repeated coin tossing of an (unfair) coin produces 100 heads up and 120 tails up. Find a maximum likelihood estimate for the probability that a coin toss will result in heads up.